


# Half-Precision Scalar Support in Kokkos and Kokkos Kernels: An Engineering Study and Experience Report



Evan Harvey, Reed Milewicz, Christian Trott,  
Luc Berger-Vergiat, Siva Rajamanickam

Research Software Engineers in eScience Workshop (RSE-eScience-2022)

10 October 2022




EXASCALE  
COMPUTING  
PROJECT



In our paper, we present a two-part study on the development of a performance portability library feature to support science and engineering applications.

- ❑ An **engineering study** on the technical implementation of the feature.
- ❑ An **experience report**, from an RSE perspective, on the challenges and lessons learned.






# Engineering Study


# What is Kokkos?



- ❑ A C++ programming model and software library ecosystem for performance portability
  - Implemented as a template library on top of CUDA, OpenMP, HPX, ...
  - Aims to be descriptive not prescriptive
  - Aligns with developments in the C++ standard
  - Replaces usage of CUDA, OpenMP, HIP, etc.
- ❑ Expanding solution for common needs of modern science/engineering codes
  - Math libraries based on Kokkos
  - Tools which enable insight into Kokkos
- ❑ Open source and widely used across a range of institutions and disciplines
  - Maintained and developed at <https://github.com/kokkos>



Performance portability  
from **laptops** to **clusters** to **supercomputers**

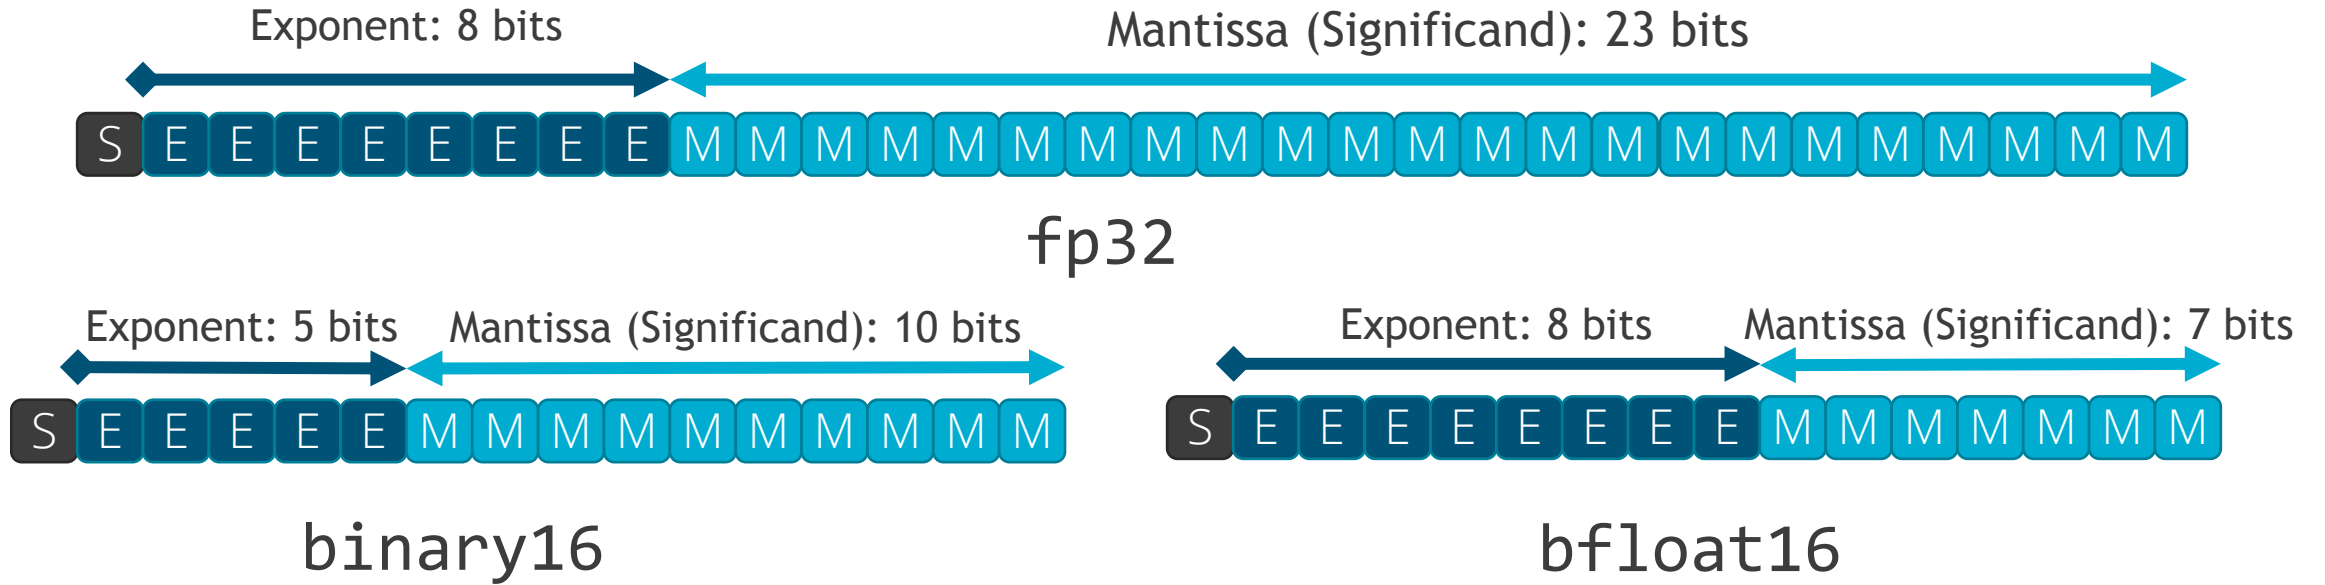


And many types of **hardware**





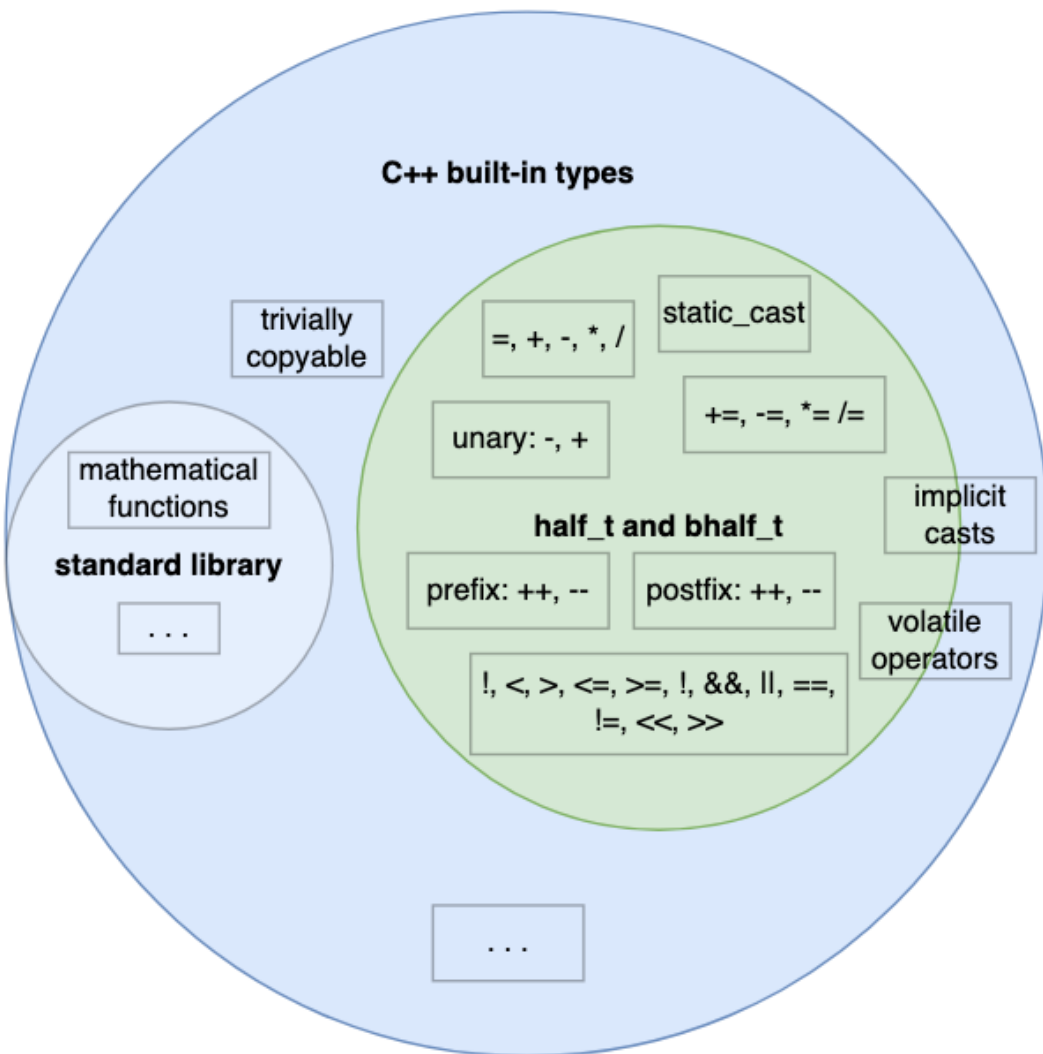
# Adding Half-Precision Floating Point Support to Kokkos



Computational science and machine learning researchers are increasingly interested in utilizing half-precision to optimize and scale their algorithms, could benefit from the addition of half-precision support.

- ❑ A 16-bit floating point encoding
- ❑ `float` encapsulates `fp32`
- ❑ `Kokkos::Experimental::half_t` encapsulates `binary16`
- ❑ `Kokkos::Experimental::bhalf_t` encapsulates `bfloat16`





□ `half_t` is either an alias to *float* or a C++ class

□ `half_t` acts like float via:

- casting wrappers with forward declarations
- operator overloading with compile-time branches

□ Volatile operations

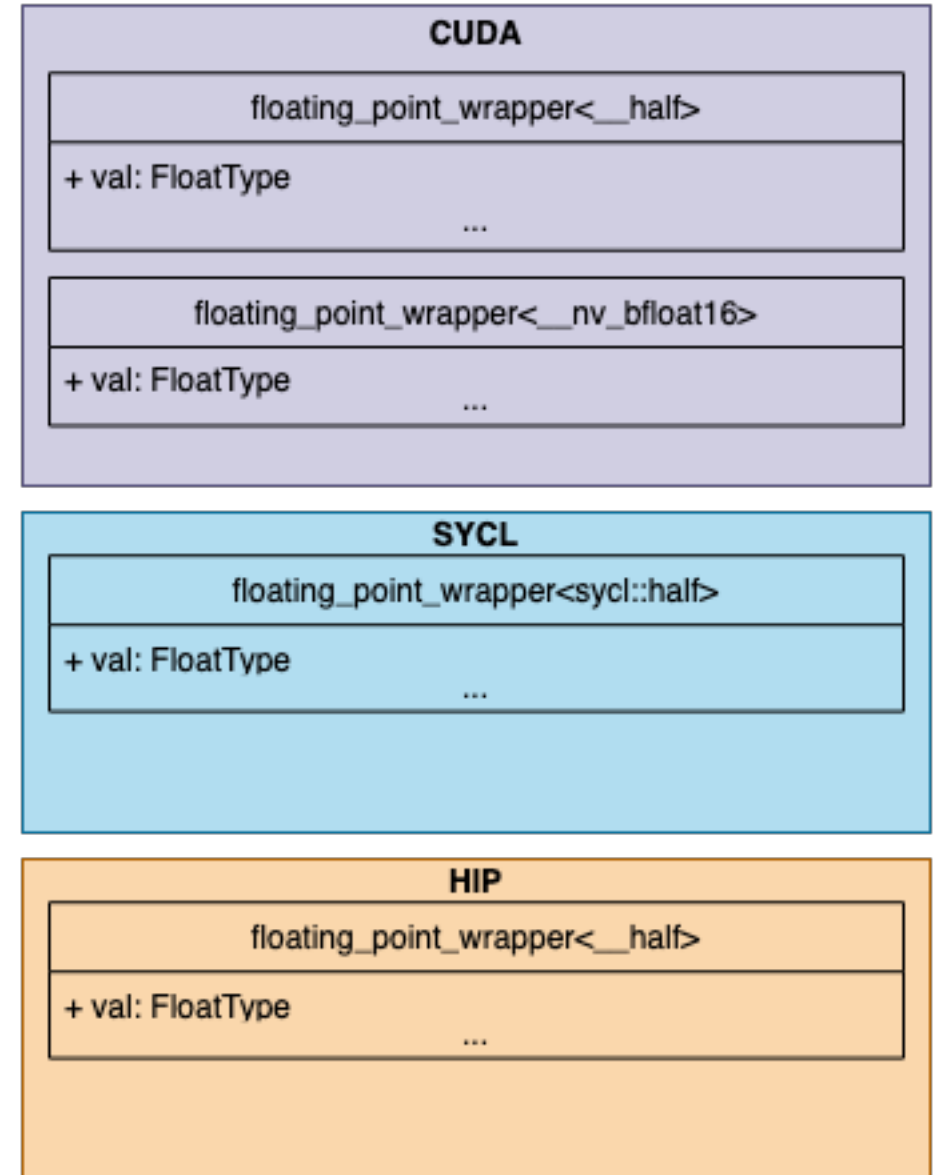
□ Mixed precision:

- `T op half_t`
- `half_t op T`





- Uses the same code as `half_t` except for:
  - Underlying data-type encodes bfloat16 via template argument
  - Casting wrappers are overloaded to call bfloat16 intrinsics

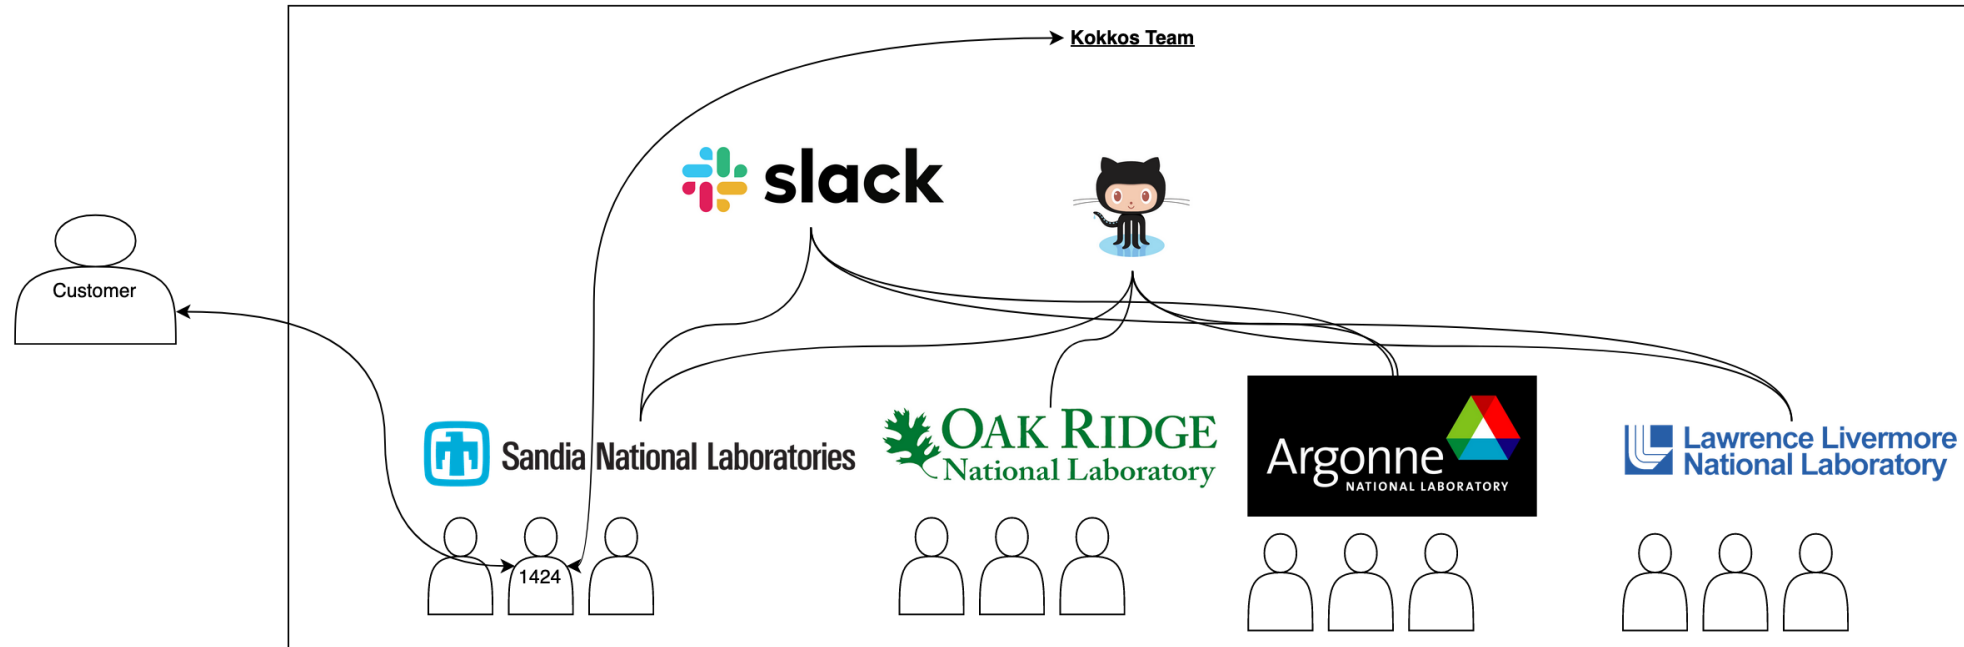




# Experience Report



## Lesson Learned: Stay Engaged with Real Users and Their Needs



Feature development for scientific software libraries should be grounded in the **needs of real users**. Proactively identify prospective stakeholders and engage with them frequently to **gather requirements**.






Be **intentional in the choice of development methodology**, and consider both your individual needs as a developer and those of your customers – different tasks may require different approaches.



## Lesson Learned: Pay Down Technical Debt Early and Often



It is important to **pay down technical debt by refactoring early and often**. Of particular note when developing scientific software libraries, latent technical debt can emerge in public interfaces and, once in place, is persistent and hard to remove.





As an RSE, **know your tools**. Case in point, modern programming languages have powerful and flexible features, but they can also be a source of complexity that must be managed. Knowing what language features to use and when is a key part of good software craftsmanship.



# Acknowledgments



- For questions and comments, feel free to reach out to me at [eharvey@sandia.gov](mailto:eharvey@sandia.gov)
- I would like to thank my mentors
  - Reed Milewicz
  - Christian Trott
  - Siva Rajamanickam
- Our sponsor
  - Exascale Computing Project
- Lastly, I would like to thank my manager

